



# Vivekanand Education Society's

## Institute of Technology

An Autonomous Institute Affiliated to University of Mumbai,, Approved by AICTE & Recognized by Govt. of Maharashtra  
Hashu Advani Memorial Complex, Collector Colony, Chembur East, Mumbai - 400074.

### Department of Information Technology

A.Y. 2024-25

## Advance DevOps Lab

### Experiment 03

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Roll No.	42
Name	NAIKWADI YASH SHIVDAS
Class	D15B
Subject	Advance DevOps Lab
LO Mapped	LO1: To understand the fundamentals of Cloud Computing and be fully proficient with Cloud based DevOps solution deployment options to meet your business requirements.  LO2: To deploy single and multiple container applications and manage application deployments with rollouts in Kubernetes
Grade:	

**Aim :** To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

## **Theory :**

### **Introduction to Kubernetes**

Kubernetes, originally developed by Google and now an open-source project under the Cloud Native Computing Foundation (CNCF), is the leading container orchestration platform. It is designed to automate the deployment, scaling, and management of containerized applications, offering developers and operations teams a robust toolset to handle the complexities introduced by container-based microservices architectures.

### **Container-Based Microservices Architecture**

With the rise of containerization, companies have moved towards microservices architectures, allowing them to deploy applications more efficiently and scale individual services independently. Containers encapsulate an application and its dependencies, ensuring consistency across development, testing, and production environments. However, the large-scale adoption of containers has also introduced new challenges, particularly in terms of managing and orchestrating hundreds or thousands of containers across multiple environments.

### **The Role of Kubernetes**

Kubernetes addresses these challenges by providing a unified platform for managing containerized workloads. It abstracts the underlying infrastructure, enabling developers to focus on building applications while Kubernetes manages the complexities of deployment and scaling.

### **Key features of Kubernetes include:**

- **Resource Management:** Kubernetes ensures that applications do not exceed resource limits set by the administrator, helping to prevent any single application from monopolizing system resources.
- **Load Balancing:** Kubernetes automatically distributes network traffic across the various instances of a service, ensuring high availability and reliability.
- **Self-Healing:** Kubernetes can detect when an application is not functioning correctly and automatically restart it or move it to a healthy node in the cluster.
- **Automated Rollouts and Rollbacks:** Kubernetes can seamlessly roll out new versions of applications and automatically roll back to a previous version in case of failure.
- **Scaling:** Kubernetes can automatically scale applications up or down based on traffic and resource utilization.
- **Cluster Management:** Kubernetes simplifies the process of adding or removing nodes from a cluster, automatically redistributing workloads as necessary.

### **Kubernetes Cluster Architecture**

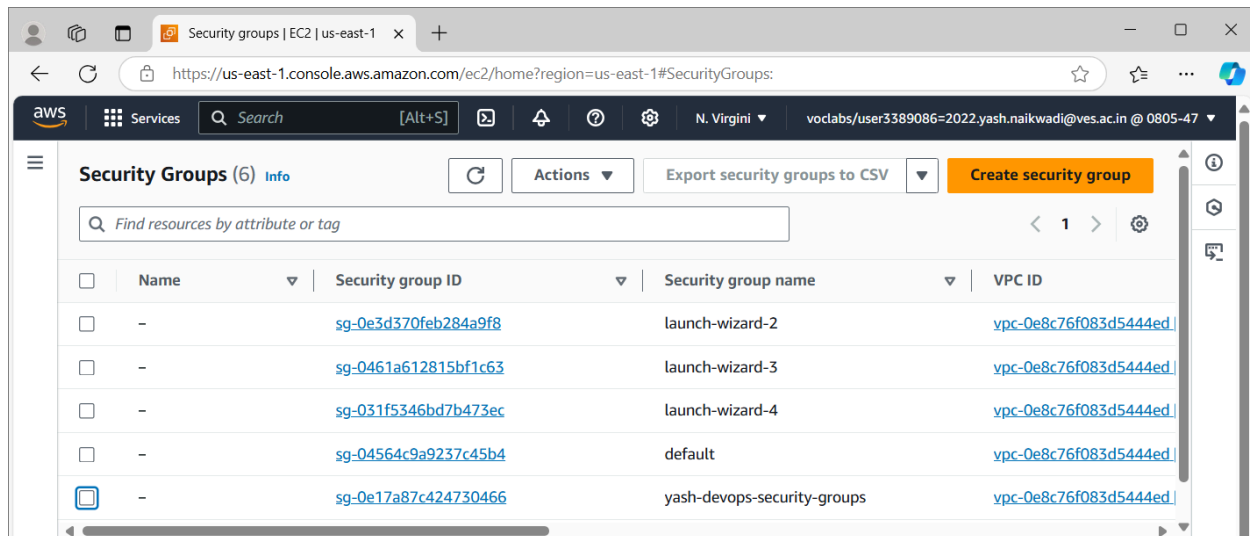
A Kubernetes cluster typically consists of a control plane (master) and worker nodes:

- **Control Plane (Master Node):** This manages the cluster and contains components like the API server, scheduler, and controller manager. The master node is responsible for maintaining the desired state of the cluster, such as which applications are running and the number of replicas.

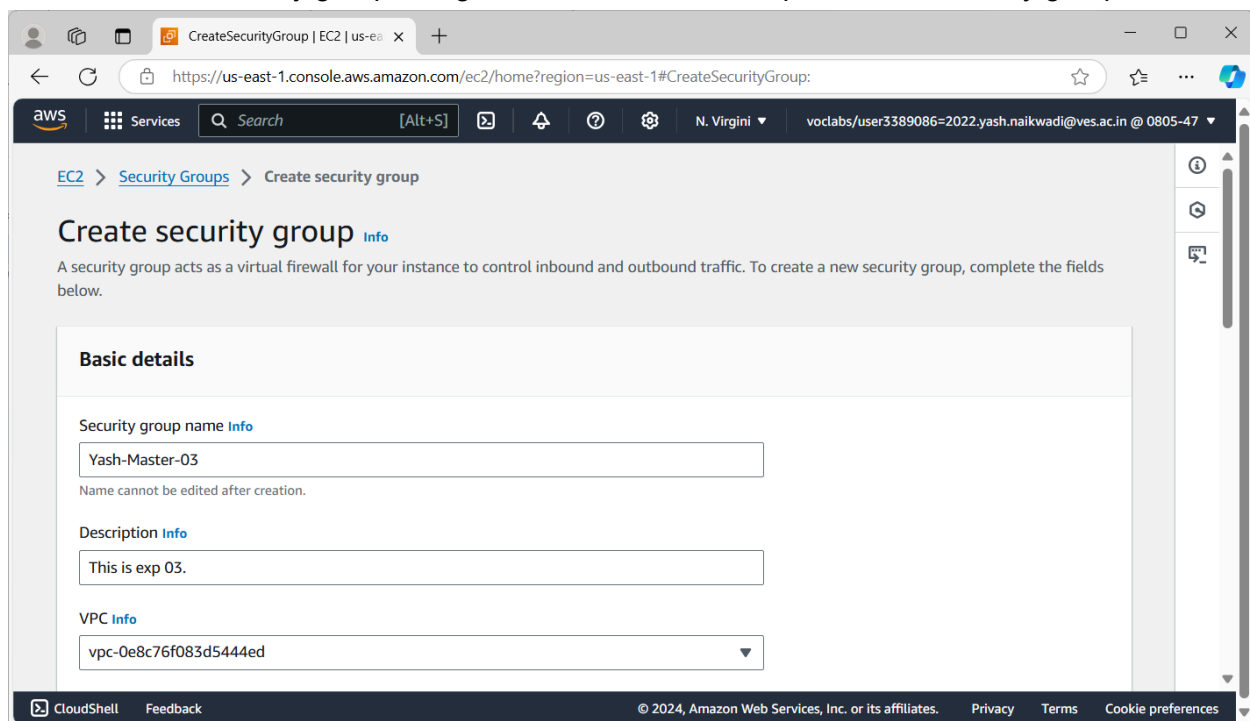
Worker Nodes: These nodes run the applications. Each worker node contains components like the Kubelet, which communicates with the control plane, and a container runtime (e.g., Docker) for running the containers.

## Create Security Groups for Instances

Log in to your **AWS account** and navigate to Security Groups (in EC2).



Click on create security group and give the name and description to the security group.



Add the following inbound rules to the security group.

The screenshot displays two pages from the AWS Management Console. The top page, titled 'Inbound rules', shows a list of rules for a security group. The bottom page, titled 'Edit outbound rules', shows a list of outbound rules for the same security group.

Type	Protocol	Port range	Source	Description - optional
HTTP	TCP	80	Anywhere...	
All traffic	All	All	Anywhere...	
Custom TCP	TCP	6443	Anywhere...	
Custom TCP	TCP	10251	Anywhere...	
Custom TCP	TCP	10250	Anywhere...	
All TCP	TCP	0 - 65535	Anywhere...	
Custom TCP	TCP	10252	Anywhere...	
SSH	TCP	22	Anywhere...	

Security group rule ID	Type	Protocol	Port range	Destination	Description - optional
sgr-058399e9b4f2d00ce	HTTPS	TCP	443	Custom	
sgr-0a51ed727dc6898e7	All traffic	All	All	Custom	
sgr-00eacb132a39f6a24	HTTP	TCP	80	Custom	

Create a security group.

The screenshot shows a notification in the AWS Management Console stating: 'Security group (sg-08fe4a3f3fb167e48 | Yash-Master-03) was created successfully'. Below the notification, the breadcrumb path is 'EC2 > Security Groups > sg-08fe4a3f3fb167e48 - Yash-Master-03'. The main heading is 'sg-08fe4a3f3fb167e48 - Yash-Master-03'.

Similarly, create one more security group (for Node).

CreateSecurityGroup | EC2 | us-east-1

https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#CreateSecurityGroup:

EC2 > Security Groups > Create security group

### Create security group [Info](#)

A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.

#### Basic details

Security group name [Info](#)

Name cannot be edited after creation.

Description [Info](#)

VPC [Info](#)

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Add the inbound rules for Node security group.

CreateSecurityGroup | EC2 | us-east-1

https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#CreateSecurityGroup:

EC2 > Security Groups > Create security group

### Create security group [Info](#)

A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.

#### Basic details

Security group name [Info](#)

Name cannot be edited after creation.

Description [Info](#)

VPC [Info](#)

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

SecurityGroup | EC2 | us-east-1

https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#SecurityGroup:groupId=sg-0f7d7b23277ad83b9

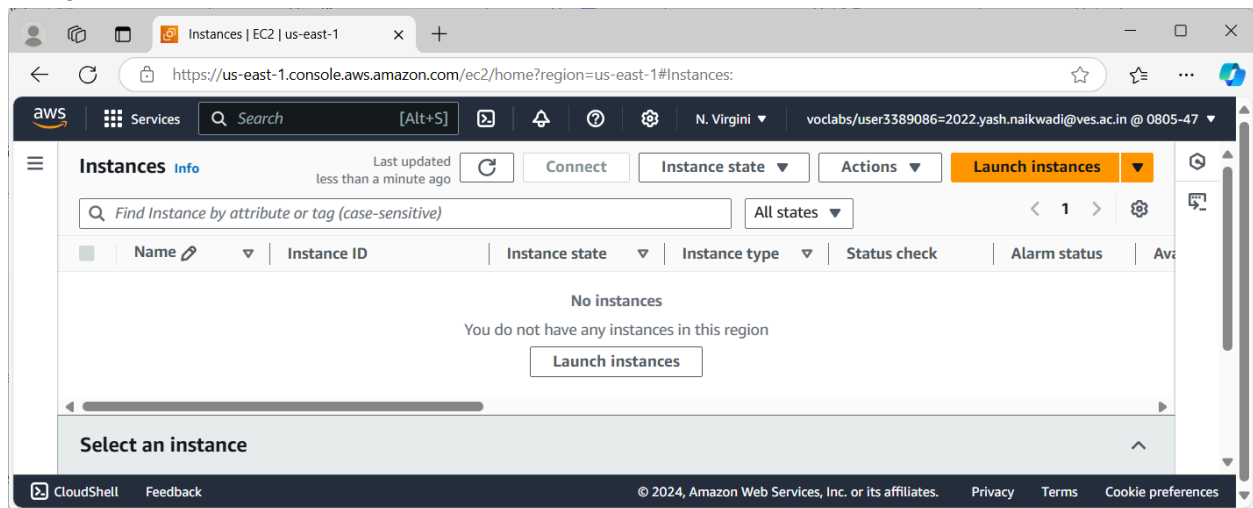
EC2 > Security Groups > sg-0f7d7b23277ad83b9 - Yash-Node-03

### sg-0f7d7b23277ad83b9 - Yash-Node-03

Actions

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

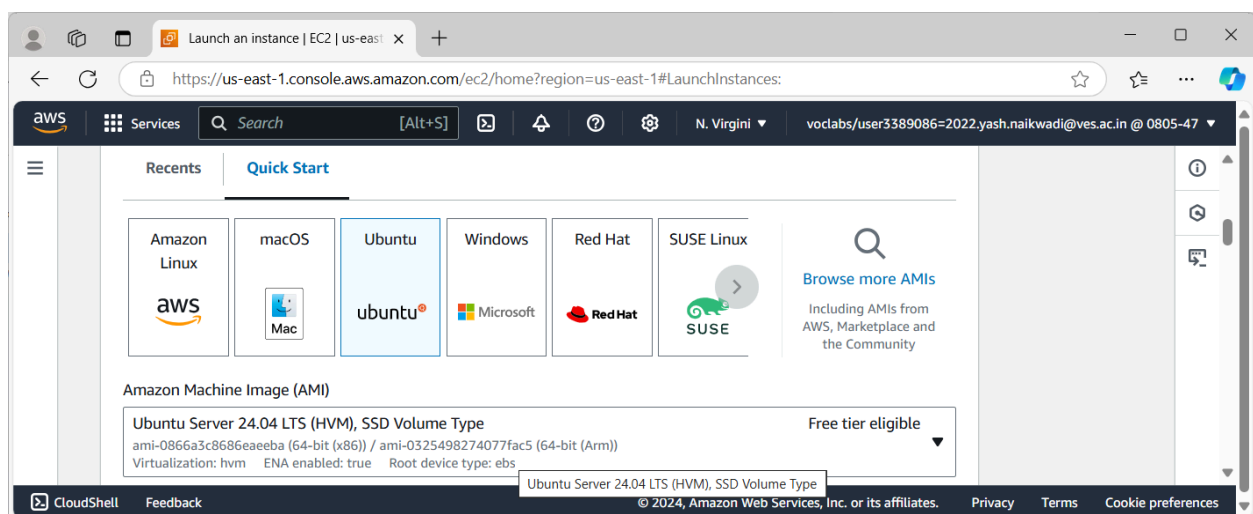
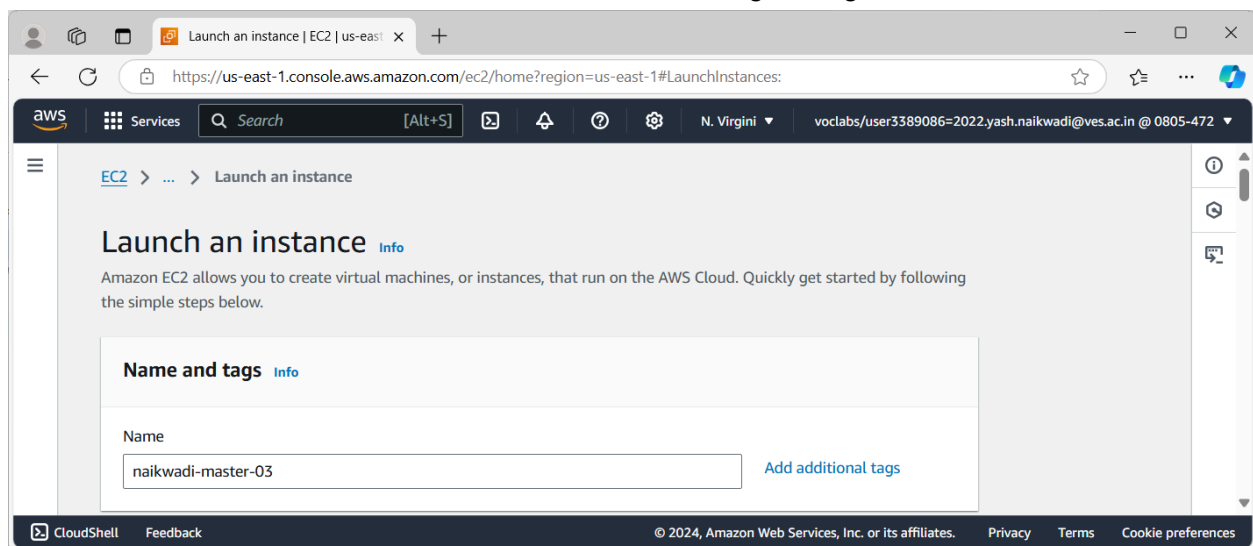
## Navigate to EC2.



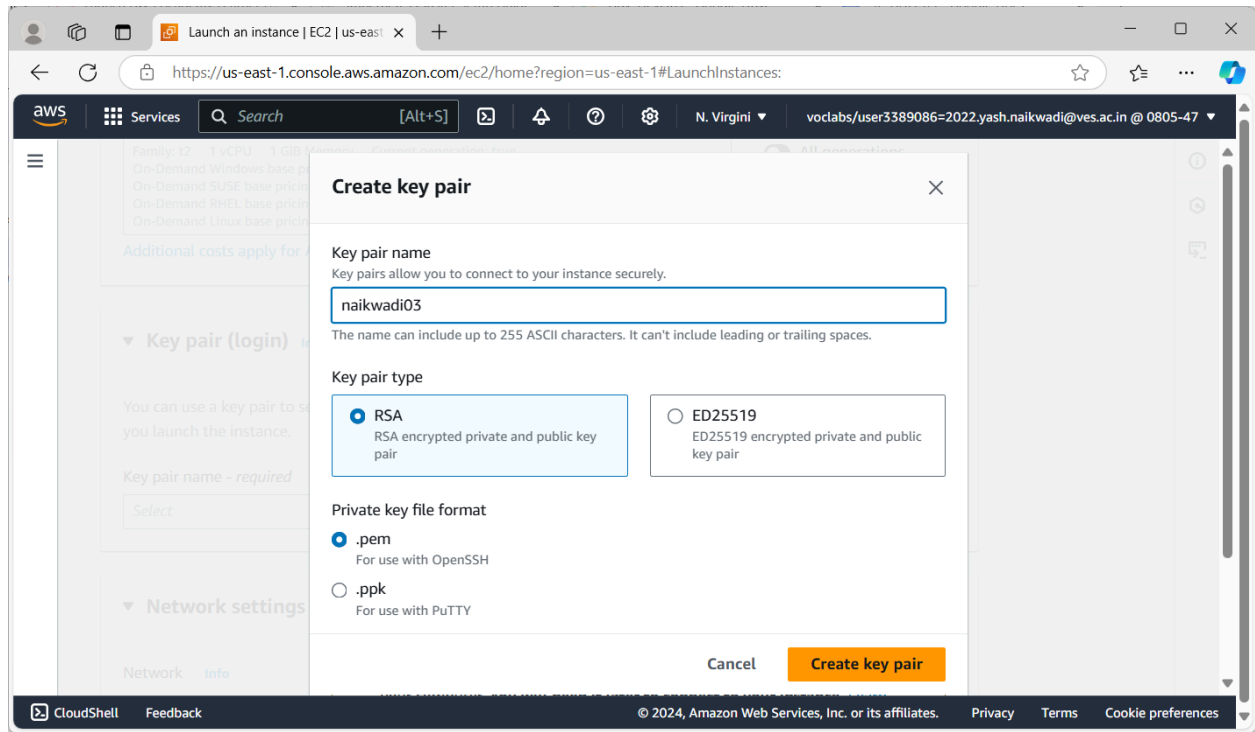
## Launch EC2 Instances

### MASTER

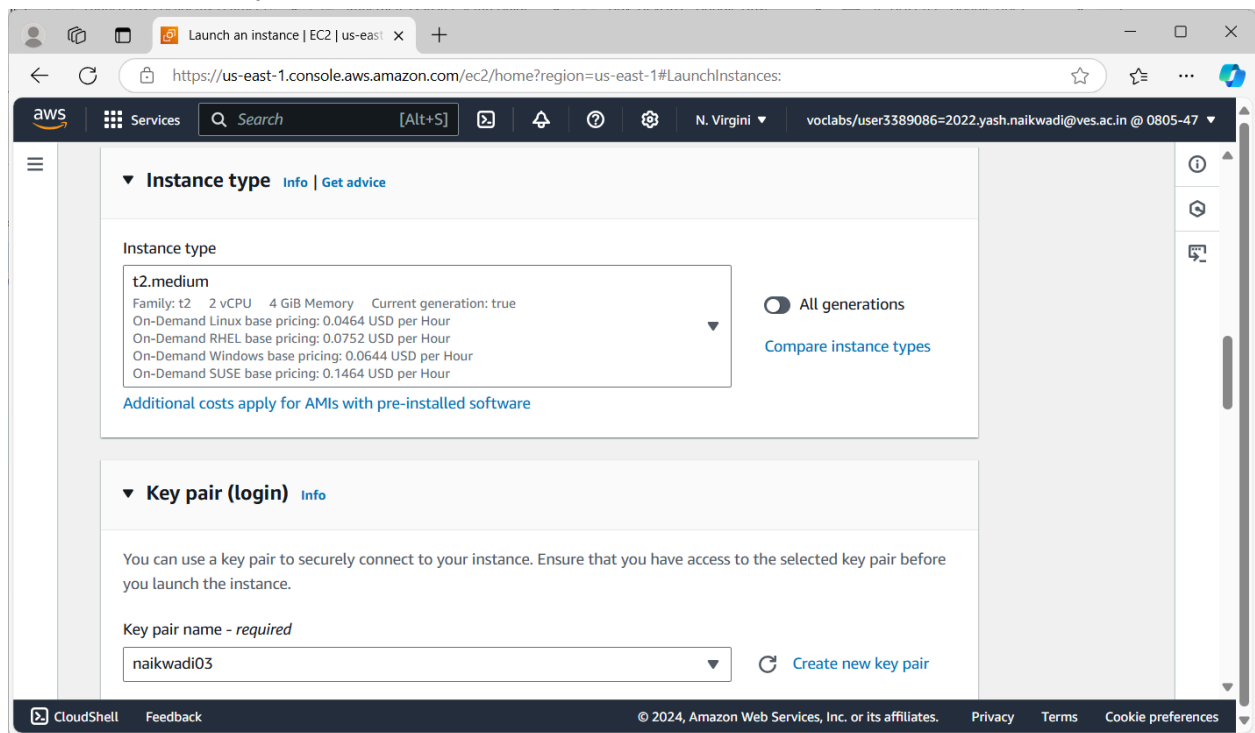
Launch a new EC2 instance for master with the following settings.



Generate the .pem file (or key-value pair) once while configuring the master ec2.

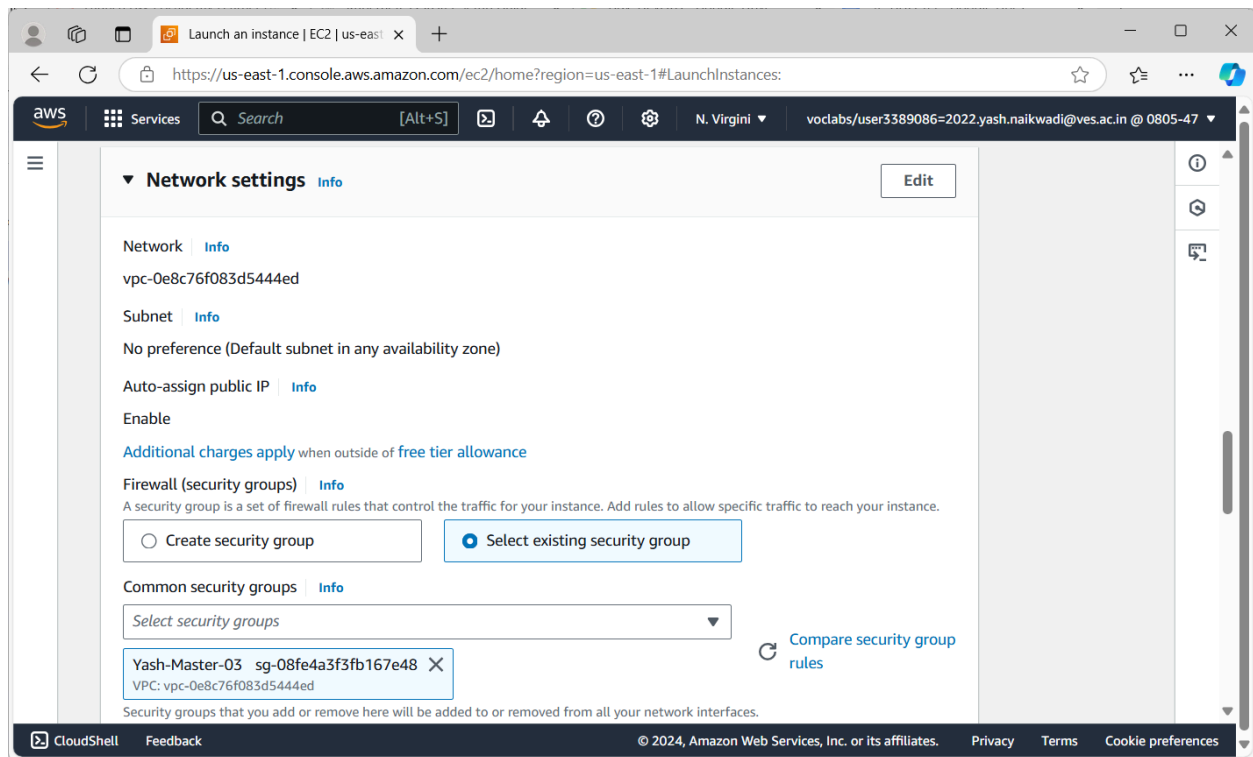


Keep the instance type as t2.medium

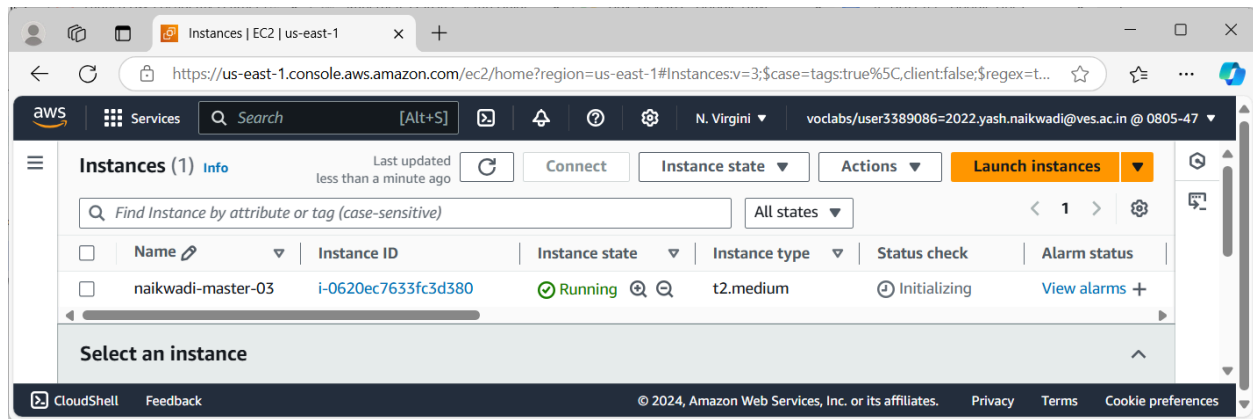


Select the existing security group under the network setting tab in which you have to select the group of master which was made in the beginning of the experiment.



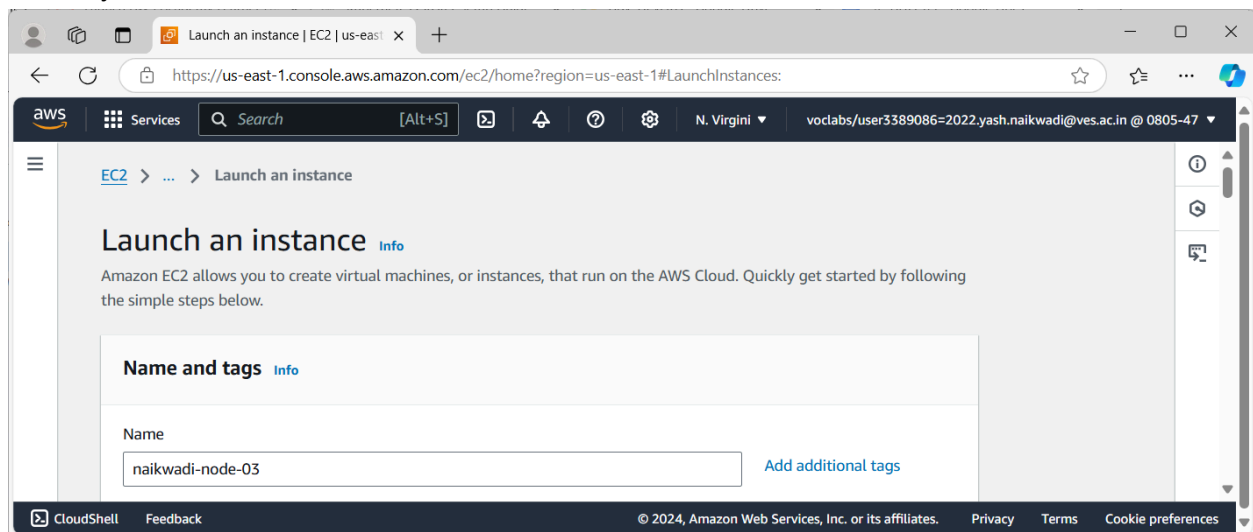


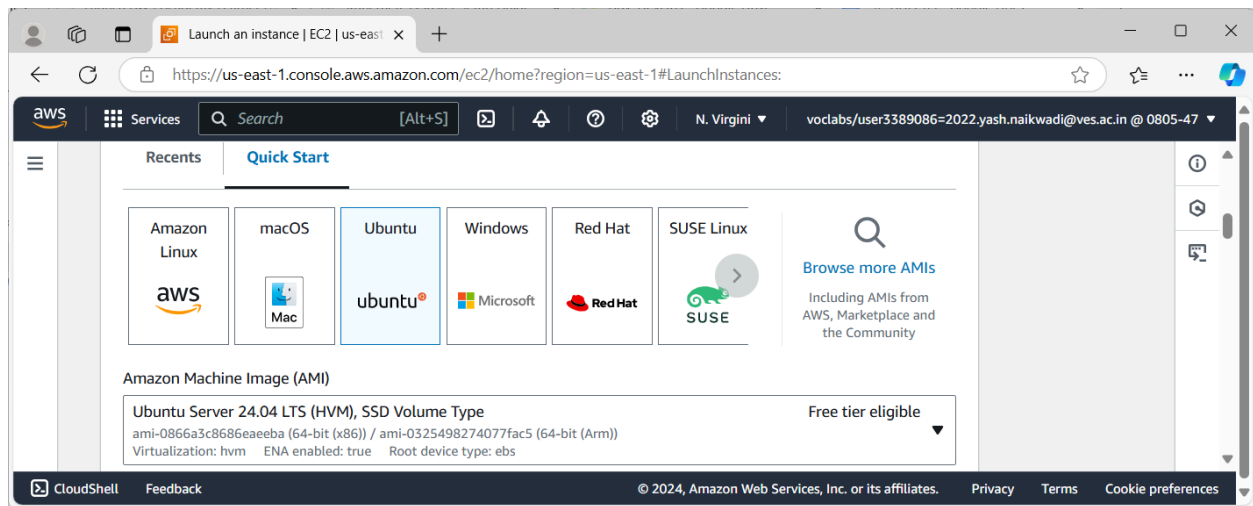
Successfully launched the ec2 instance.



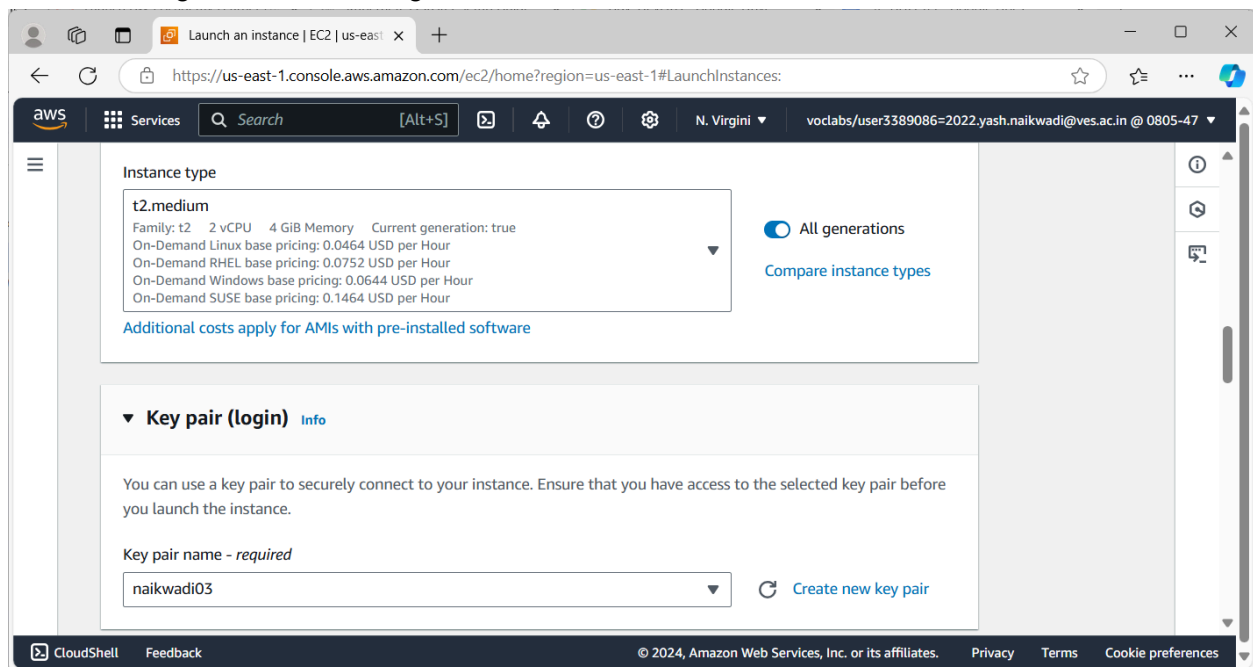
## NODE

Similarly, launch 2 new ec2 instances for nodes.

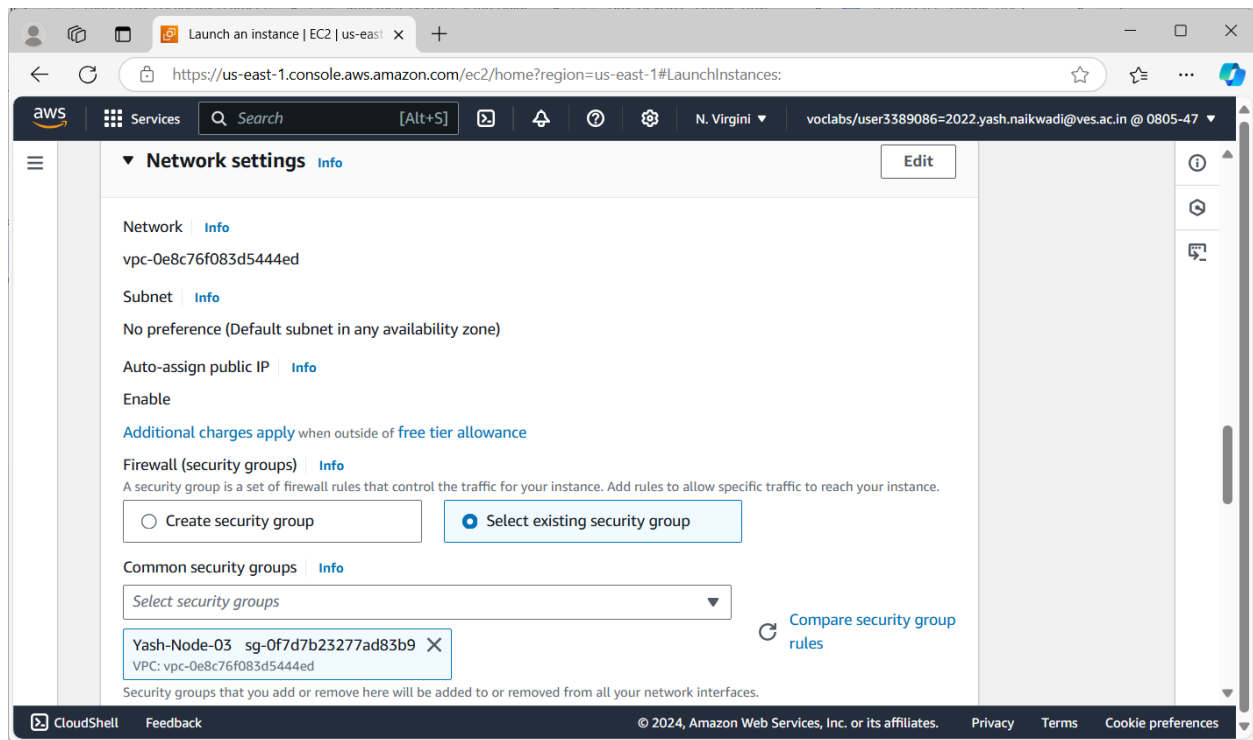




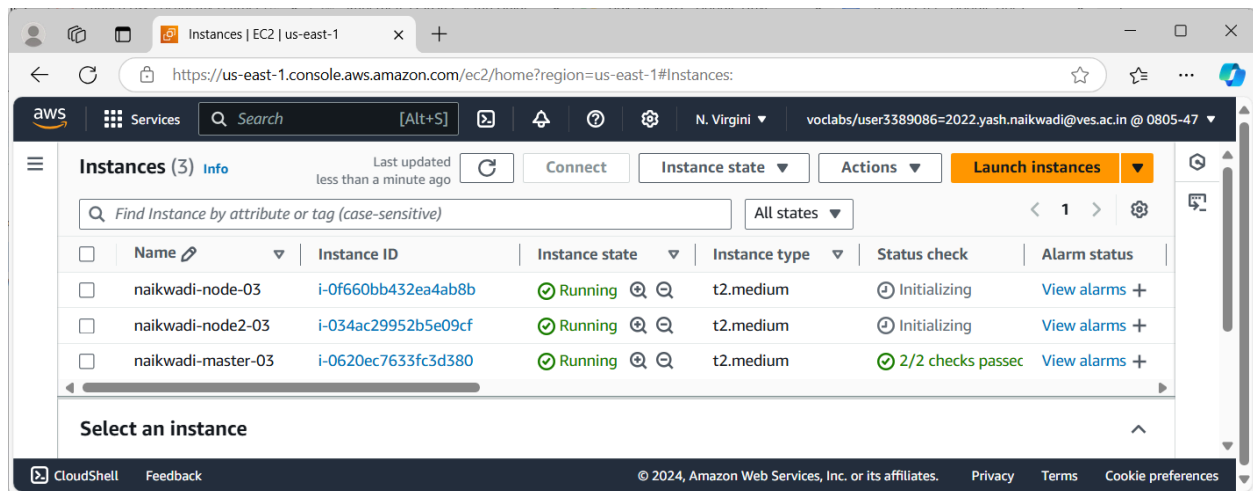
Keep the instance type of node as t2.medium and select the same key-value pair which was created during the master configuration.



Select the existing security group which was made for Node in the beginning of the experiment.

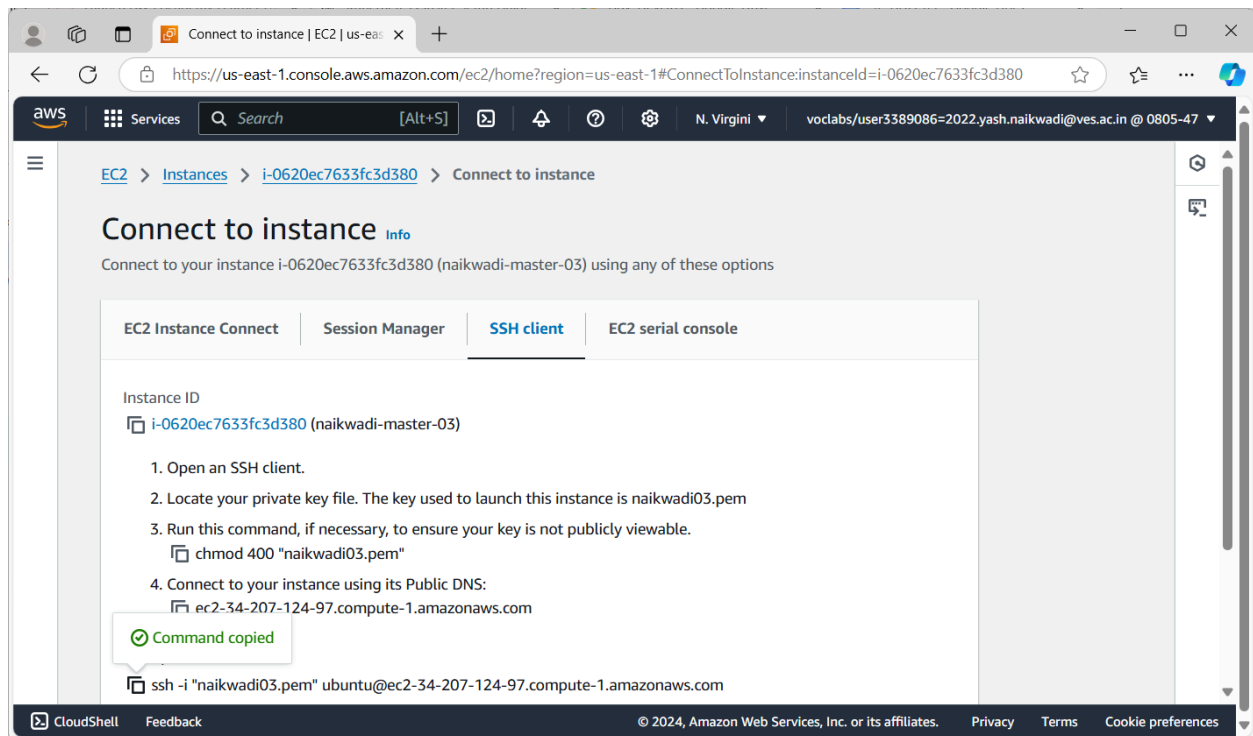


Successfully, created the 3 ec2 instances (1 for master & 2 for nodes).

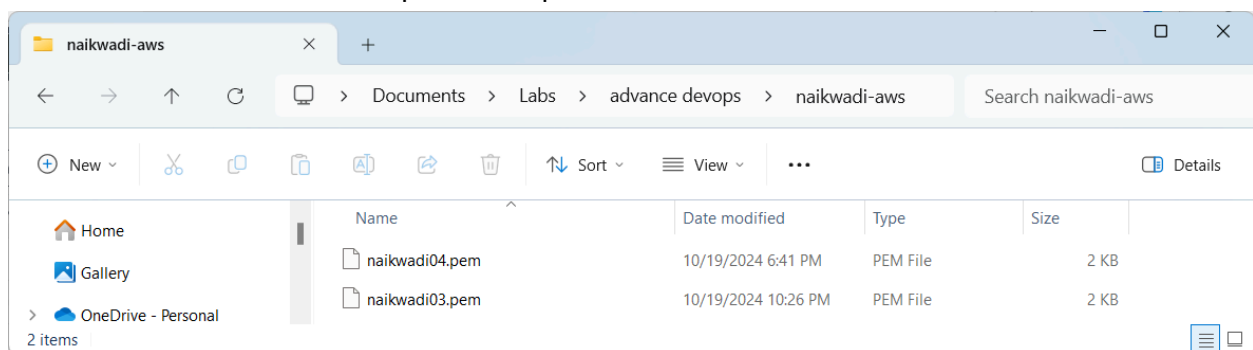


### Connect to the Instances

After launching the instances, go to the **EC2 dashboard**, select each instance, and click **Connect** to get the SSH command.

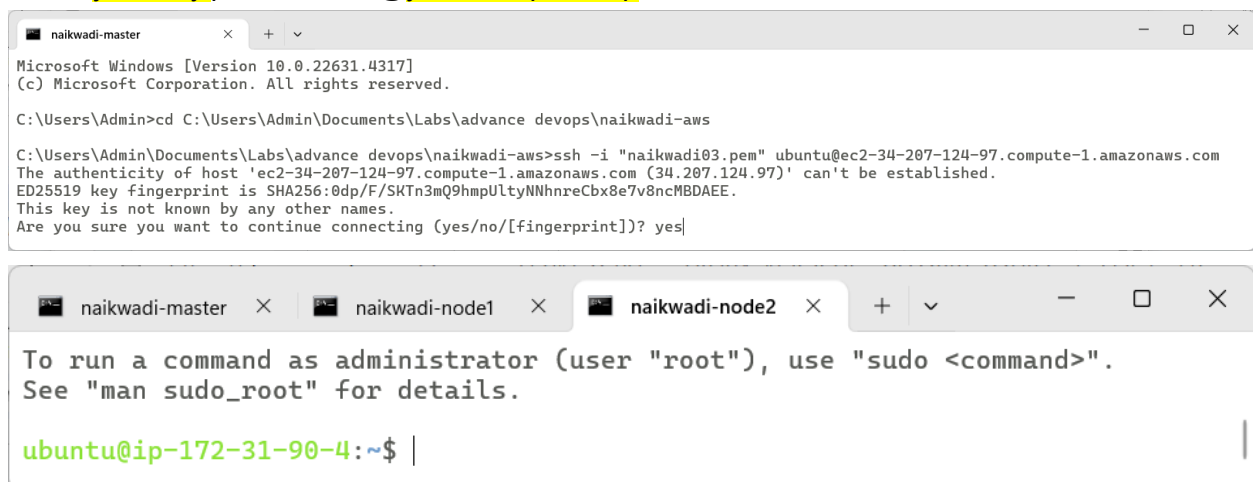


Locate the folder in which the .pem file is placed.



Open your terminal, navigate to the folder where your **.pem key** is stored, and connect to each instance using:

`ssh -i "your-key.pem" ubuntu@your-ec2-public-ip`



### Install Docker on All Instances (Master and Nodes)

On each instance (Master, Node 1, Node 2), run the following commands to install Docker:

# Add Docker's official GPG key

`curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -`

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-87-71:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).
OK
ubuntu@ip-172-31-87-71:~$ |
```

### # Add Docker's APT repository

```
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
$(lsb_release -cs) stable"
```

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-88-198:~$ sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
Repository: 'deb [arch=amd64] https://download.docker.com/linux/ubuntu noble stable'
Description:
Archive for codename: noble components: stable
More info: https://download.docker.com/linux/ubuntu
Adding repository.
Press [ENTER] to continue or Ctrl-c to cancel.
```

### # Update the package index

```
sudo apt-get update
```

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-88-198:~$ sudo apt-get update
Ign:1 https://download.docker.com/linux/ubuntu noble InRelease
Ign:2 http://security.ubuntu.com/ubuntu noble-security InRelease
Ign:1 https://download.docker.com/linux/ubuntu noble InRelease
Ign:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Ign:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Ign:5 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
```

### # Install Docker CE (Community Edition)

```
sudo apt-get install -y docker-ce
```

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-90-4:~$ sudo apt-get install -y docker-ce
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  containerd.io docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
```

### # Configure Docker to use systemd as the cgroup driver

```
sudo mkdir -p /etc/docker
```

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-88-198:~$ sudo mkdir -p /etc/docker
ubuntu@ip-172-31-88-198:~$
```

```
cat <<EOF | sudo tee /etc/docker/daemon.json
```

```
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
```

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-90-4:~$ cat <<EOF | sudo tee /etc/docker/daemon.json
> {
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
ubuntu@ip-172-31-90-4:~$ |
```

# Enable and start Docker

sudo systemctl enable docker

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-88-198:~$ sudo systemctl enable docker
Synchronizing state of docker.service with SysV service script with /usr/lib/sys
temd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
ubuntu@ip-172-31-88-198:~$ |
```

sudo systemctl daemon-reload

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-90-4:~$ sudo systemctl daemon-reload
ubuntu@ip-172-31-90-4:~$ |
```

sudo systemctl restart docker

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-88-198:~$ sudo systemctl restart docker
ubuntu@ip-172-31-88-198:~$ |
```

### Install Kubernetes on All Instances

Run the below command to install Kubernetes.

curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-87-71:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Relea
se.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
File '/etc/apt/keyrings/kubernetes-apt-keyring.gpg' exists. Overwrite? (y/N) Y
ubuntu@ip-172-31-87-71:~$ |
```

echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]

https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-87-71:~$ echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pk
gs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/de
b/ /
ubuntu@ip-172-31-87-71:~$ |
```

sudo apt-get update

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-87-71:~$ sudo apt-get update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 https://download.docker.com/linux/ubuntu noble InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
```

sudo apt-get install -y kubelet kubeadm kubectl

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-87-71:~$ sudo apt-get install -y kubelet kubeadm kubectl
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  contrack cri-tools kubernetes-cni
The following NEW packages will be installed:
  contrack cri-tools kubeadm kubectl kubelet kubernetes-cni
The following held packages will be changed:
  kubeadm kubectl kubelet
0 upgraded, 6 newly installed, 0 to remove and 25 not upgraded.
E: Held packages were changed and -y was used without --allow-change-held-packages.
ubuntu@ip-172-31-87-71:~$
```

sudo apt-mark hold kubelet kubeadm kubectl

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-87-71:~$ sudo apt-mark hold kubelet kubeadm kubectl
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
ubuntu@ip-172-31-87-71:~$
```

sudo systemctl enable --now kubelet

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-88-198:~$ sudo systemctl enable --now kubelet
ubuntu@ip-172-31-88-198:~$
```

sudo apt-get install -y containerd

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-87-71:~$ sudo apt-get install -y containerd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
```

sudo mkdir -p /etc/containerd

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + v - □ x
ubuntu@ip-172-31-88-198:~$ sudo mkdir -p /etc/containerd
ubuntu@ip-172-31-88-198:~$
```

sudo containerd config default | sudo tee /etc/containerd/config.toml



```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + - □ x
ubuntu@ip-172-31-88-198:~$ sudo containerd config default | sudo tee /etc/containerd/config.toml
disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2
```

sudo systemctl restart containerd

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + - □ x
ubuntu@ip-172-31-88-198:~$ sudo systemctl restart containerd
ubuntu@ip-172-31-88-198:~$
```

sudo systemctl enable containerd

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + - □ x
ubuntu@ip-172-31-88-198:~$ sudo systemctl enable containerd
ubuntu@ip-172-31-88-198:~$
```

sudo systemctl status containerd

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + - □ x
ubuntu@ip-172-31-90-4:~$ sudo systemctl status containerd
● containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; enabled; preset: ➤
   Active: active (running) since Sun 2024-10-20 03:37:45 UTC; 40s ago
     Docs: https://containerd.io
   Main PID: 7312 (containerd)
    Tasks: 7
   Memory: 13.4M (peak: 13.9M)
      CPU: 158ms
   CGroup: /system.slice/containerd.service
           └─7312 /usr/bin/containerd

Oct 20 03:37:45 ip-172-31-90-4 containerd[7312]: time="2024-10-20T03:37:45.2530512"
lines 1-12...skipping...
```

sudo apt-get install -y socat

```
naikwadi-master x naikwadi-node1 x naikwadi-node2 x + - □ x
ubuntu@ip-172-31-90-4:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 25 not upgraded.
```

### Initialize the Kubernetes Master Node

Initialize the Kubecuster .Now Perform this Command only for Master.

sudo kubeadm init --pod-network-cidr=10.244.0.0/16



```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  X
ubuntu@ip-172-31-88-198:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.1
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your i
nternet connection

```

Once the initialization is complete, you will see a command that looks like this:

(**Copy and save** this command for later; you will need it to join the worker nodes to the master.)

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  X
kubeadm join 172.31.88.198:6443 --token sfxllk.6ggcci840yv22h18 \
--discovery-token-ca-cert-hash sha256:1a08766f3ac51776ba839909a5e22b7
9a96f4b353260439519866eca40a50203
ubuntu@ip-172-31-88-198:~$

```

Next, configure kubectl to interact with the cluster:

`mkdir -p $HOME/.kube`

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  X
ubuntu@ip-172-31-88-198:~$ mkdir -p $HOME/.kube
ubuntu@ip-172-31-88-198:~$

```

`sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  X
ubuntu@ip-172-31-88-198:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
ubuntu@ip-172-31-88-198:~$

```

`sudo chown $(id -u):$(id -g) $HOME/.kube/config`

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  X
ubuntu@ip-172-31-88-198:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@ip-172-31-88-198:~$

```

Now Run the command

`kubectl get nodes`

to see the nodes before executing the Join command on nodes.

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  X
ubuntu@ip-172-31-88-198:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE     VERSION
ip-172-31-88-198    NotReady control-plane 5m36s    v1.31.1
ubuntu@ip-172-31-88-198:~$

```

### Join Worker Nodes to the Cluster

On **Node 1** and **Node 2**, run the `kubeadm join` command you saved from the master node initialization. It should look like this:

`sudo kubeadm join <Master-IP>:6443 --token <token> --discovery-token-ca-cert-hash <hash>`

(In my case, command is as follows:

`sudo kubeadm join 172.31.88.198:6443 --token sfxllk.6ggcci840yv22h18 \`

`--discovery-token-ca-cert-hash`

`sha256:1a08766f3ac51776ba839909a5e22b79a96f4b353260439519866eca40a50203)`

```

naikwadi-master x  naivekadi-node1 x  naivekadi-node2 x  +  -  □  ×
ubuntu@ip-172-31-87-71:~$ sudo kubeadm join 172.31.88.198:6443 --token sfxllk.6ggcci840yv22h18 --discovery-token-ca-cert-hash sha256:1a08766f3ac51776ba839909a5e22b79a96f4b353260439519866eca40a50203
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 501.460744ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-87-71:~$ |

```

```

naikwadi-master x  naivekadi-node1 x  naivekadi-node2 x  +  -  □  ×
ubuntu@ip-172-31-90-4:~$ sudo kubeadm join 172.31.88.198:6443 --token sfxllk.6ggcci840yv22h18 --discovery-token-ca-cert-hash sha256:1a08766f3ac51776ba839909a5e22b79a96f4b353260439519866eca40a50203
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.001793491s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-90-4:~$ |

```

### Verify the Cluster

On the master node, verify that the nodes have joined the cluster:

kubectl get nodes

```

naikwadi-master x  naivekadi-node1 x  naivekadi-node2 x  +  -  □  ×
ubuntu@ip-172-31-88-198:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-87-71     NotReady <none>   84s   v1.31.1
ip-172-31-88-198    NotReady control-plane 13m   v1.31.1
ip-172-31-90-4      NotReady <none>   21s   v1.31.1
ubuntu@ip-172-31-88-198:~$ |

```

### Install a Network Plugin (Calico)

Since Status is NotReady we have to add a network plugin. And also we have to give the name to the nodes.

kubectl apply -f <https://docs.projectcalico.org/manifests/calico.yaml>

```

naikwadi-master x  naivekadi-node1 x  naivekadi-node2 x  +  -  □  ×
ubuntu@ip-172-31-88-198:~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created

```

sudo systemctl status kubelet

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  ×

ubuntu@ip-172-31-88-198:~$ sudo systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; preset: enabled)
   Drop-In: /usr/lib/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Sun 2024-10-20 03:41:35 UTC; 15min ago
     Docs: https://kubernetes.io/docs/
   Lines 1-6...skipping...
● kubelet.service - kubelet: The Kubernetes Node Agent

```

Now Run command (we can see Status is ready).

kubectl get nodes -o wide

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  ×

ubuntu@ip-172-31-88-198:~$ kubectl get nodes -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-87-71     Ready     <none>    5m7s  v1.31.1   172.31.87.71  <none>        Ubuntu 24.04.1 LTS   6.8.0-1016-aws   containerd://1.7.12
ip-172-31-88-198    Ready     control-plane 17m   v1.31.1   172.31.88.198 <none>        Ubuntu 24.04.1 LTS   6.8.0-1016-aws   containerd://1.7.12
ip-172-31-90-4      Ready     <none>    4m4s  v1.31.1   172.31.90.4   <none>        Ubuntu 24.04.1 LTS   6.8.0-1016-aws   containerd://1.7.12
ubuntu@ip-172-31-88-198:~$ |

```

### Rename/Label the Nodes

kubectl label node ip-172-31-87-71 kubernetes.io/role=Node1

kubectl label node ip-172-31-90-4 kubernetes.io/role=Node2

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  ×

ubuntu@ip-172-31-88-198:~$ kubectl label node ip-172-31-87-71 kubernetes.io/role=Node1
node/ip-172-31-87-71 labeled
ubuntu@ip-172-31-88-198:~$ kubectl label node ip-172-31-90-4 kubernetes.io/role=Node2
node/ip-172-31-90-4 labeled
ubuntu@ip-172-31-88-198:~$ |

```

### Final Check

Run (again to confirm all nodes are in the **Ready** state and properly labeled).

kubectl get nodes

```

naikwadi-master x  naikwadi-node1 x  naikwadi-node2 x  +  -  □  ×

ubuntu@ip-172-31-88-198:~$ kubectl get nodes -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-87-71     Ready     Node1     10m   v1.31.1   172.31.87.71  <none>        Ubuntu 24.04.1 LTS   6.8.0-1016-aws   containerd://1.7.12
ip-172-31-88-198    Ready     control-plane 22m   v1.31.1   172.31.88.198 <none>        Ubuntu 24.04.1 LTS   6.8.0-1016-aws   containerd://1.7.12
ip-172-31-90-4      Ready     Node2     9m12s v1.31.1   172.31.90.4   <none>        Ubuntu 24.04.1 LTS   6.8.0-1016-aws   containerd://1.7.12
ubuntu@ip-172-31-88-198:~$ |

```

Or run

kubectl get nodes

```

naikwadi- x  naikwadi- x  naikwadi- x  +  -  □  ×

ubuntu@ip-172-31-88-198:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-87-71     Ready     Node1     10m   v1.31.1
ip-172-31-88-198    Ready     control-plane 22m   v1.31.1
ip-172-31-90-4      Ready     Node2     9m48s v1.31.1
ubuntu@ip-172-31-88-198:~$ |

```

### Conclusion :

In this experiment, we successfully set up a Kubernetes cluster on AWS EC2 instances, understanding how Kubernetes manages and orchestrates containerized applications across multiple nodes. This hands-on experience demonstrated Kubernetes' ability to automate deployment, scaling, and resource management in a cloud environment.